

Decision-Dependent Risk Minimization in Geometrically Decaying Dynamic Environments

- **Date:** 2022
- **Link:** [paper](#)
- **Authors:**
 - [Ray, Mitas](#)
 - [Ratliff, Lillian J](#)
 - [Drusvyatskiy, Dmitriy](#)
 - [Maryam, Fazel](#)
- **Cites:**
- **Cited by:**
- **Keywords:** [#performative-prediction](#) [#decision-dependent-risk-min](#)
- **Collections:**
- **Status:** [#completed](#)

0 - Abstract

- Problem description:
 - Perform expected loss minimization.
 - The given data distribution is dependent on the decision-maker's action.
 - The distribution evolves dynamically in time according to a geometric decay process.
- Proposed algorithms
 - Two information settings:
 - Decision-maker has a first order gradient oracle.
 - Decision-maker has a loss function oracle.
 - Key principle:
 - The decision-maker repeatedly deploys a fixed decision repeatedly over the length of an epoch, allowing the environment to sufficiently mix before updating the decision.
- Results:
 - Iteration complexity is shown to match existing rates for first and zero order stochastic gradient methods up to logarithmic factors.
 - Algorithms are evaluated on SF Park dynamic pricing pilot study.
 - Algorithm is shown to improve target occupancy and reduce overall parking rates.

1 - Introduction

- Supervised machine learning algorithms are trained on past data under the assumption that the past data is representative of the future.
- This is not the case in examples where the output of the algorithm changes the environment:
 - Online labor markets
 - Predictive policing
 - On-street parking
 - Vehicle sharing markets
- *Performative prediction* (aka decision-dependent risk minimization)- solution to this problem, models the data distribution as being *decision dependent* thereby accounting for feedback induced distributional shift.
- Past work has only focused on static environments; however, most real-world environments are not static.
 - i.e. data distribution changes dynamically in time as the strategic actors adjust their behavior based on the algorithm's change.
 - e.g. if a city changes its street parking pricing, it may take time for drivers to modify their behavior in response. After some time, an equilibrium will be achieved.
- Algorithmic framework:
 - Start of the epoch:
 - Decision-maker selects a decision to deploy repeatedly for the duration of the epoch.
 - During the epoch:
 - The distribution evolves toward a fixed point distribution for the given decision.
 - At the end of the epoch:
 - Decision-maker is updated using a first-order or zeroth-order oracle.
(Depending on the information setting - gradient or loss function access).

1.1 - Contributions

- Iteration complexity guarantees:

- **Zero Order Oracle** (Algorithm 1)
 - Sample complexity is $\tilde{O}(d^2/\epsilon^2)$.
 - This matches the optimal rate.
- **First Order Oracle** (Algorithm 2)
 - Sample complexity is $\tilde{O}(1/\epsilon)$.
 - This matches the known optimal rate.
- This is achieved by bounding the error between the expected gradient at the fixed point distribution and the stochastic gradient at time t .
- Algorithm was shown to produce positive results on the SF Park pilot study.
 - Higher occupancy rates for parking spots.
 - Lower total parking fees.

1.2 - Related Work

1.2.1 - Dynamic Decision-Dependent Optimization

- *Recourse* - the decision-maker is able to make a secondary decision after some information has been revealed.
- Dynamic decision-dependent optimization has been extensively studied in the case of recourse.
- This is unlike the setting considered here.

1.2.2 - Reinforcement Learning

- RL is a closely related problem:
 - Agent makes decisions in a dynamic environment that responds to the agent's actions.
- Important difference:
 - Our goal
 - Find the action which optimizes the decision-dependent expected risk at the fixed point distribution.
 - RL's goal
 - Find the policy which is a state-dependent distribution over actions given an accumulated cost.
- In this sense, our setting can be viewed as a special case of general RL.

1.2.3 - Performative prediction

- Naïve strategy - retrain the model after using heuristics to determine when there is sufficient distribution shift.
- None of the works consider dynamic environments.

2 - Preliminaries

- Metric space notation:
 - Let \mathcal{Z} be a metric space with metric $d(\cdot, \cdot)$
 - Let $\mathbb{P}(\mathcal{Z})$ denote the set of Radon probability measures ν on \mathcal{Z} with finite first moment.
 - *Note to self* - I need to understand this part better.

- Decision-maker seeks to solve:

$$\min_{x \in \mathcal{X}} \mathcal{L}(x)$$

with the expected loss function

$$\mathcal{L}(x) = E_{z \sim \mathcal{D}(x)} [\ell(x, z)]$$

- Definitions:
 - \mathcal{X} - the decision space in \mathbb{R}^d
 - x - the decision made by the decision-maker.
 - z - the response from the environment.
 - $\mathcal{D}(x)$ - the probability distribution from which z is sampled, depends on x and t .
- Let x^* denote an optimal decision vector.
- Probability measure:
 - **Main challenge** - the response, characterized as a random variable, z depends not only on the decision x_t , but also explicitly on the time step t .
 - Let z be governed by the distribution p_t .
 - p_t is generated by the process - $p_{t+1} = \mathcal{T}(p_t, x_t)$ where

$$\mathcal{T}(p, x) = \lambda p + (1 - \lambda) \mathcal{D}(x)$$

is a geometric decay rate.

- Interpretation - at each time step t , $(1 - \lambda)$ fraction of the population becomes aware of the decision x .
- Alternative interpretation - the environment has a memory that is captured by the distribution, as time progresses past observations decay at a rate of λ .

- See the paper for a list of assumptions for ℓ and $\mathcal{D}(x)$.

3 - Algorithms and Sample Complexity Analysis

- Recall, the decision-maker holds fixed the decision for n time steps before observing the environment.

3.1 - Zero Order Stochastic Gradient Method

- This is the most general information setting in which we assume the decision-maker does not have access to $\mathcal{D}(x)$.
 - In practice, this is likely to be true.

- Stochastic gradient method:
 - Let $\delta > 0$ - this 'wiggles' the decision point around during the epoch.

- At each epoch t , the algorithm...
 - Samples v_t from a d -dimensional unit sphere.
 - Queries the environment for n_t iterations with $x_t + \delta v_t$.

- At the end of the epoch...
 - The loss oracle reveals $\ell(x_t + \delta v_t, z_t)$ where

$$z_t \sim \lambda^{n_t} p_{t-1} + (1 - \lambda^{n_t}) \mathcal{D}(x_t + \delta v_t)$$

Note - this is the accumulated environmental response distribution after n_t repeated selections of $x_t + \delta v_t$.

- The decision-maker performs the update rule:

$$x_{t+1} = \text{proj}_{(1-\delta)\mathcal{X}}(x_t - \eta \hat{g}_t)$$

where

$$\hat{g}_t = \frac{d}{\delta} \ell(x_t + \delta v_t, z_t) v_t$$

is the one-point gradient estimate of the expected loss at p_t .

- *Note to self:* I'm not sure why the projection is taken here? Paper says its to ensure x in the next iteration is in the feasible set.
- Smoothed expected risk:

$$\mathcal{L}^\delta(x) = E_{v \sim B} \left[E_{z \sim \mathcal{D}(x + \delta v)} [\ell(x + \delta v, z)] \right]$$

This is strongly convex.

- See the paper for the sample complexity proof.

3.2 - First Order Stochastic Gradient Method

- In this information setting, we assume the decision-maker has access to a parametric description of $\mathcal{D}(x)$.

- Expected loss:

$$\mathcal{L}_t(x) = E_{z \sim p_t} \ell(x_t, z)$$

- Differentiating (under a mild smoothness assumption):

$$\nabla \mathcal{L}_t(x) = E_{z \sim p_t} [\nabla_x \ell(x, z) + (1 - \lambda^n) A^T \nabla_z \ell(x, z)]$$

- Therefore, given a decision x , the decision-maker can always sample a response $z \sim p_t$ and form the unbiased estimator of $E_{z \sim p_t} [\hat{g}_t] = \nabla \mathcal{L}_t(x)$:

$$\hat{g}_t = \nabla \ell(x_t, z) = \nabla_x \ell(x, z) + (1 - \lambda^n) A^T \nabla_z \ell(x, z)$$

- Algorithm:
 - For each round t :
 - Decision-maker queries the environment with x_t for n steps.
 - Resulting in the distribution, $p_t = \lambda^n p_{t-1} + (1 - \lambda^n) \mathcal{D}(x_t)$.
 - The gradient oracle reveals \hat{g}_t
 - And, the update rule is applied $x_{t+1} = \text{proj}_{\mathcal{X}}(x_t - \eta_t \hat{g}_t)$.

- See the paper for the sample complexity proof.

4 - Numerical Experiments

- Target metric - between 60 and 80% occupancy rates.

- Rate structure:
 - Operational hours are split into three distinct rate periods.
 - Rates are adjusted on a block-by-block basis.
 - Rates are adjusted based on current occupancy.
 - Only weekday, non-special event parking is considered.

- Problem setup:
 - Decision x is the changes in price from the nominal price relative to SF Park's initial rates.

- Fixed point distribution z is the curb occupancies.
- Dynamic decision-dependent loss:

$$E_{z \sim p_t} [\ell(x, z)] = E [\|z - 0.7\|^2 + \frac{\nu}{2} \|x\|^2]$$

where ν is the regularization parameter.

- Data distribution is defined as follows:

$$z \sim \mathcal{D}(x) \iff z = \zeta + Ax$$

- ζ - follows the distribution p_0 which is sampled from the data at the beginning of the pilot study.
- A is a proxy for price elasticity which is estimated via linear fit between occupancy and price.

4.1 - Comparing Performative Optimum to SF Park

- Run the algorithms on a representative block.
- Chosen parameters:
 - $A \approx -0.157$ - a \$1 increase in rates results in an approx. 15% decrease in occupancy.
 - $\lambda \approx 0.959$ - decay rate, see appendix for calculation.
 - SF Park initial rate for the block is \$3.
 - $\mathcal{X} = [-3, 5]$ - decision space - min rate is \$0; max rate is \$8.
 - $\nu = 1e - 3$ - regularization param.
 - Time step and epoch length are varied, in units of weeks.
- Analysis:
 - Choice of (n, T) for Algorithm 1:
 - For the sample curb, $n = 8$ appears optimal .
 - Generally lower n and higher T yielded better results.
 - Different optimal n were observed for different curbs.
 - Suggesting a non-uniform price update schedule.
 - Choice of (n, T) for Algorithm 2:
 - Opposite is true, generally higher n and lower T yielded better results.
 - This is due to the variance in the query direction for Algorithm 1.

4.2 - Redistributing Parking Demand

- Examine a case with a high demand block (Hawthorne St 0) surrounded by lower demand blocks.
- Algorithm is able to reduce the occupancy at the high demand block from 90% to 70% and redistribute this to the lower demand blocks.
- This result is much better than SF Park.
- Interestingly, the algorithm does not perform well on a block where the unconstrained equilibrium price is higher than the max \$8, set my SF Park.

5 - Discussion and Future Directions

- Future work:
 - Experiment with period dynamics.
 - May be possible to exploit additional structure on $\mathcal{D}(x)$ to improve sample complexity.