

No-Regret Learning in Bilateral Trade via Global Budget Balance

- **Date:** 2024-06-24
- [Link](#)
- **Authors:**
 - [Bernasconi, Martino](#)
 - [Castiglioni, Matteo](#)
 - [Celli, Andrea](#)
 - [Fusco, Federico](#)
- **Cites:**
- **Cited by:**
- **Keywords:** [#bilateral-trade](#) [#online-bilateral-trade](#) [#global-budget-balancing](#) [#bandits-with-knapsacks](#)
- **Collections:**
- **Status:** [#in-progress](#)

0. Abstract

- Problem description:
 - *Bilateral trade* - problem of facilitating trade between a buyer and a seller, each of whom contain a private valuation of the item for sale.
 - *Online bilateral trade* - online learning problem in which a new buyer and seller are introduced at each time step, and the learner has no prior knowledge of their preferences.
- **Problem** - there does not exist a no-regret algorithm that is budget balanced at each timestep.
 - This is a proven result.
- **Solution** - instead of requiring the learner to budget balance at each timestep, enforce a balanced budget over the entire time horizon.
- **Result** - first no-regret algorithms for adversarial bilateral trade.
 - Full feedback model - show $\tilde{O}(\sqrt{T})$ regret against the best fixed prices in hindsight.
 - This is an optimal bound (up to poly-logarithmic terms).
 - One-bit feedback model - show $\tilde{O}(T^{3/4})$ regret upper bound and $\Omega(T^{5/7})$ lower bound.
- An alternative, stronger benchmark is introduced, inspired by the literature on bandits with knapsacks.